

AN7914 Week 05 Python

February 26, 2024

1 Week 5 Python

1.1 Introduction to Pandas

Pandas is a package built on top of **NumPy**, and provides an efficient implementation of a **DataFrame**. **DataFrames** are essentially multidimensional *arrays* with attached row and column labels, and often with heterogeneous types and/or missing data.

1.2 Installing Pandas

Installation of Pandas on your system requires NumPy to be installed. Details on this installation can be found in the [Pandas](#) documentation. Once Pandas is installed, you can import it and check the version:

```
import pandas
pd.__version__
```

We will however use an alias to call pandas. So when importing we do the following:

```
[1]: import pandas as pd
```

In the code above we imported **pandas** under the alias **pd**. Now let's check the version again, but this we will use the alias.

```
[2]: pd.__version__
```

```
[2]: '1.5.2'
```

1.3 Creating data

There are two core objects in pandas: the **DataFrame** and the **Series**.

1.4 DataFrame

A **DataFrame** is a table. It contains an array of individual entries, each of which has a certain value. Each entry corresponds to a *row* (or record) and a *column*.

For example, consider the following simple DataFrame:

```
[3]: pd.DataFrame({'Yup': [50, 21, 32], 'Nope': [131, 2, 200]})
```

```
[3]:    Yup  Nope
      0   50  131
      1   21   2
      2   32  200
```

`DataFrame` entries are not limited to integers. For instance, here's a `DataFrame` whose values are strings:

```
[4]: pd.DataFrame({'Bob': ['I liked it.', 'It was awful.'], 'Sue': ['Pretty good.',
    ↪ 'Bland.']}])
```

```
[4]:           Bob           Sue
0    I liked it.  Pretty good.
1  It was awful.    Bland.
```

We are using the `pd.DataFrame()` constructor to generate these `DataFrame` objects. The syntax for declaring a new one is a **dictionary** whose keys are the *column* names (Bob and Sue in this example), and whose *values* are a list of entries.

The dictionary-list constructor assigns values to the column labels, but just uses an ascending count from 0 (0, 1, 2, 3, ...) for the row labels. Sometimes this is OK, but oftentimes we will want to assign these labels ourselves.

```
[5]: pd.DataFrame({'Bob': ['I liked it.', 'It was awful.'],
    'Sue': ['Pretty good.', 'Bland.']],
    index=['Product A', 'Product B'])
```

```
[5]:           Bob           Sue
Product A    I liked it.  Pretty good.
Product B  It was awful.    Bland.
```

1.5 Series

A **Series**, by contrast, is a sequence of data values. If a `DataFrame` is a table, a **Series** is a **list**. And in fact you can create one with nothing more than a list:

```
[6]: pd.Series([1,2,3,4])
```

```
[6]: 0    1
      1    2
      2    3
      3    4
      dtype: int64
```

A **Series** is, in essence, a single **column** of a `DataFrame`. So you can assign row labels to the **Series** the same way as before, using an **index** parameter. However, a **Series** does not have a column name, it only has one overall **name**:

1.6 Importing data sets

You will need to use and import data sets from the internet or from your hard-drive. So if you want to import a `csv` file you will need to use `pd.read_csv()` command. The argument in the command could be the location where the file is stored in your computer or it could be a file store in the internet as show below.

```
[7]: employee=pd.read_csv('/Users/sakibanwar/Downloads/updated_employee_dataset.csv')
```

We use `pd.read_csv()` to read a file stored in `'/Users/sakibanwar/Downloads/updated_employee_dataset.csv'`. Then we store this dataset as a `dataframe` in `employee`.

Now lets take a look at the dataset. We can simply type `employee` the name of the `dataframe`. It will usually not going to show all columns and rows.

```
[8]: employee
```

```
[8]:
```

	Employee_ID	Department	Years_of_Experience	Full_Time	Performance_Score	\
0	1	Finance	1.82	False	Good	
1	2	IT	2.09	False	Poor	
2	3	IT	13.24	False	Excellent	
3	4	Finance	1.55	False	Excellent	
4	5	Marketing	4.16	True	Average	
5	6	Finance	7.42	True	Excellent	
6	7	Marketing	10.42	True	Good	
7	8	Marketing	12.90	True	Good	
8	9	Marketing	10.00	True	Poor	
9	10	HR	18.74	False	Average	
10	11	Finance	6.17	False	Poor	
11	12	IT	11.26	True	Average	
12	13	HR	19.10	True	Good	
13	14	IT	6.46	True	Excellent	
14	15	Finance	14.47	True	Good	
15	16	IT	7.40	True	Average	
16	17	IT	4.56	False	Poor	
17	18	HR	12.48	True	Poor	
18	19	HR	3.77	True	Excellent	
19	20	HR	4.86	True	Excellent	
20	21	IT	15.91	False	Average	
21	22	IT	15.97	True	Poor	
22	23	Marketing	8.06	False	Excellent	
23	24	Finance	18.83	True	Average	
24	25	Marketing	11.52	True	Excellent	
25	26	Finance	8.55	True	Excellent	
26	27	HR	10.74	True	Good	
27	28	HR	5.13	False	Average	
28	29	Finance	8.09	False	Good	
29	30	Finance	8.82	True	Excellent	
30	31	IT	19.08	True	Good	

31	32	Finance	2.08	False	Poor
32	33	HR	3.65	False	Poor
33	34	HR	15.79	True	Poor
34	35	Finance	15.13	True	Poor
35	36	IT	7.27	False	Average
36	37	IT	13.26	False	Poor
37	38	IT	19.42	False	Good
38	39	Finance	16.00	False	Good
39	40	HR	11.64	True	Excellent
40	41	Finance	1.12	True	Average
41	42	Marketing	14.48	False	Good
42	43	Marketing	15.61	True	Poor
43	44	HR	2.28	False	Average
44	45	Finance	11.35	True	Excellent
45	46	Finance	14.65	True	Good
46	47	IT	7.38	True	Average
47	48	IT	16.16	True	Excellent
48	49	IT	5.72	True	Average
49	50	IT	4.32	True	Good

	Salary	First_Name	Last_Name
0	54550.0	Michael	Davis
1	55225.0	Karen	Brown
2	83100.0	Joseph	Johnson
3	53875.0	David	Garcia
4	60400.0	Linda	Martinez
5	68550.0	Michael	Brown
6	76050.0	Charles	Moore
7	82250.0	David	Lopez
8	75000.0	David	Johnson
9	96850.0	Patricia	Martinez
10	65425.0	Linda	Brown
11	78150.0	John	Wilson
12	97750.0	Mary	Thomas
13	66150.0	Barbara	Anderson
14	86175.0	Jennifer	Davis
15	68500.0	Mary	Wilson
16	61400.0	James	Davis
17	81200.0	Barbara	Thomas
18	59425.0	Barbara	Gonzalez
19	62150.0	Thomas	Moore
20	89775.0	Elizabeth	Anderson
21	89925.0	Jessica	Gonzalez
22	70150.0	Joseph	Rodriguez
23	97075.0	Joseph	Anderson
24	78800.0	Charles	Gonzalez
25	71375.0	Barbara	Smith

26	76850.0	Karen	Miller
27	62825.0	John	Rodriguez
28	70225.0	Robert	Smith
29	72050.0	Charles	Lopez
30	97700.0	Michael	Davis
31	55200.0	William	Hernandez
32	59125.0	Michael	Jackson
33	89475.0	Sarah	Taylor
34	87825.0	Patricia	Davis
35	68175.0	Susan	Williams
36	83150.0	Sarah	Williams
37	98550.0	William	Smith
38	90000.0	Mary	Jones
39	79100.0	Karen	Martinez
40	52800.0	Joseph	Miller
41	86200.0	Michael	Rodriguez
42	89025.0	Barbara	Miller
43	55700.0	Linda	Rodriguez
44	78375.0	Joseph	Davis
45	86625.0	John	Lopez
46	68450.0	Susan	Johnson
47	90400.0	Thomas	Smith
48	64300.0	Patricia	Thomas
49	60800.0	Sarah	Jones

Now let's take a look at the first 15 rows.

```
[9]: employee.head(15)
```

```
[9]:
```

	Employee_ID	Department	Years_of_Experience	Full_Time	Performance_Score	\
0	1	Finance	1.82	False	Good	
1	2	IT	2.09	False	Poor	
2	3	IT	13.24	False	Excellent	
3	4	Finance	1.55	False	Excellent	
4	5	Marketing	4.16	True	Average	
5	6	Finance	7.42	True	Excellent	
6	7	Marketing	10.42	True	Good	
7	8	Marketing	12.90	True	Good	
8	9	Marketing	10.00	True	Poor	
9	10	HR	18.74	False	Average	
10	11	Finance	6.17	False	Poor	
11	12	IT	11.26	True	Average	
12	13	HR	19.10	True	Good	
13	14	IT	6.46	True	Excellent	
14	15	Finance	14.47	True	Good	

```
Salary First_Name Last_Name
```

0	54550.0	Michael	Davis
1	55225.0	Karen	Brown
2	83100.0	Joseph	Johnson
3	53875.0	David	Garcia
4	60400.0	Linda	Martinez
5	68550.0	Michael	Brown
6	76050.0	Charles	Moore
7	82250.0	David	Lopez
8	75000.0	David	Johnson
9	96850.0	Patricia	Martinez
10	65425.0	Linda	Brown
11	78150.0	John	Wilson
12	97750.0	Mary	Thomas
13	66150.0	Barbara	Anderson
14	86175.0	Jennifer	Davis

`employee.head(15)` gives us the first 15 rows of the dataframe. If we typed `employee.head(25)` it would show us the first 25 rows.

To see the last 10 rows we can use `employee.tail(10)`

```
[10]: employee.tail(10)
```

```
[10]:
```

	Employee_ID	Department	Years_of_Experience	Full_Time	Performance_Score	\
40	41	Finance	1.12	True	Average	
41	42	Marketing	14.48	False	Good	
42	43	Marketing	15.61	True	Poor	
43	44	HR	2.28	False	Average	
44	45	Finance	11.35	True	Excellent	
45	46	Finance	14.65	True	Good	
46	47	IT	7.38	True	Average	
47	48	IT	16.16	True	Excellent	
48	49	IT	5.72	True	Average	
49	50	IT	4.32	True	Good	

	Salary	First_Name	Last_Name
40	52800.0	Joseph	Miller
41	86200.0	Michael	Rodriguez
42	89025.0	Barbara	Miller
43	55700.0	Linda	Rodriguez
44	78375.0	Joseph	Davis
45	86625.0	John	Lopez
46	68450.0	Susan	Johnson
47	90400.0	Thomas	Smith
48	64300.0	Patricia	Thomas
49	60800.0	Sarah	Jones

If you want to know exactly how many rows and columns the dataframe has we can simply type

```
employee.shape
```

```
[11]: employee.shape
```

```
[11]: (50, 8)
```

We see that the output is (50, 8)– this means we have 50 rows and 8 columns.

If we want to see the names of the columns we use `employee.columns`

```
[12]: employee.columns
```

```
[12]: Index(['Employee_ID', 'Department', 'Years_of_Experience', 'Full_Time',  
          'Performance_Score', 'Salary', 'First_Name', 'Last_Name'],  
          dtype='object')
```

If we want to see a specific columns we can pass in a list –

`employee[['Performance_Score', 'Salary']]`. We passed in the list `['Performance_Score', 'Salary']`. This list contains the list of column names. The output is going to look a dataframe and not a series.

```
[13]: employee[['Performance_Score', 'Salary']]
```

```
[13]:
```

	Performance_Score	Salary
0	Good	54550.0
1	Poor	55225.0
2	Excellent	83100.0
3	Excellent	53875.0
4	Average	60400.0
5	Excellent	68550.0
6	Good	76050.0
7	Good	82250.0
8	Poor	75000.0
9	Average	96850.0
10	Poor	65425.0
11	Average	78150.0
12	Good	97750.0
13	Excellent	66150.0
14	Good	86175.0
15	Average	68500.0
16	Poor	61400.0
17	Poor	81200.0
18	Excellent	59425.0
19	Excellent	62150.0
20	Average	89775.0
21	Poor	89925.0
22	Excellent	70150.0
23	Average	97075.0

24	Excellent	78800.0
25	Excellent	71375.0
26	Good	76850.0
27	Average	62825.0
28	Good	70225.0
29	Excellent	72050.0
30	Good	97700.0
31	Poor	55200.0
32	Poor	59125.0
33	Poor	89475.0
34	Poor	87825.0
35	Average	68175.0
36	Poor	83150.0
37	Good	98550.0
38	Good	90000.0
39	Excellent	79100.0
40	Average	52800.0
41	Good	86200.0
42	Poor	89025.0
43	Average	55700.0
44	Excellent	78375.0
45	Good	86625.0
46	Average	68450.0
47	Excellent	90400.0
48	Average	64300.0
49	Good	60800.0

1.7 Indexing in Pandas

Two types of indexing: 1. Index-based selection– `.iloc`

2. Label-based selection– `.loc`

1.8 Index-based selection– `.iloc`

`.iloc` is about selecting data based on its numerical position. Let's use `.iloc` to select the first row in the `hotel_vienna` DataFrame.

```
[14]: employee.iloc[0]
```

```
[14]: Employee_ID      1
      Department      Finance
      Years_of_Experience  1.82
      Full_Time        False
      Performance_Score    Good
      Salary          54550.0
      First_Name      Michael
      Last_Name       Davis
```


Name: 0, dtype: object

This returned a `Series` object. If we want to get the second row. We do the following:

```
[15]: employee.iloc[1]
```

```
[15]: Employee_ID      2
      Department      IT
      Years_of_Experience  2.09
      Full_Time      False
      Performance_Score  Poor
      Salary      55225.0
      First_Name      Karen
      Last_Name      Brown
      Name: 1, dtype: object
```

This returned a `Series` object. If we want to get the third row. We do the following:

```
[16]: employee.iloc[2]
```

```
[16]: Employee_ID      3
      Department      IT
      Years_of_Experience  13.24
      Full_Time      False
      Performance_Score  Excellent
      Salary      83100.0
      First_Name      Joseph
      Last_Name      Johnson
      Name: 2, dtype: object
```

We are going by row-indexes here! You get the idea! How about looking at multiple rows at once?
We can slice rows

```
[17]: employee.iloc[0:3]
```

```
[17]:   Employee_ID  Department  Years_of_Experience  Full_Time  Performance_Score  \
0           1    Finance           1.82      False             Good
1           2         IT           2.09      False             Poor
2           3         IT          13.24      False          Excellent

      Salary  First_Name  Last_Name
0  54550.0    Michael    Davis
1  55225.0      Karen    Brown
2  83100.0     Joseph  Johnson
```

In the above we used `employee.iloc[0:3]` to slice the rows and get the rows whose index are numbered 0,1,2. Another example below:

```
[18]: employee.iloc[0:10]
```

```
[18]: Employee_ID Department Years_of_Experience Full_Time Performance_Score \
0      1      Finance      1.82      False      Good
1      2      IT      2.09      False      Poor
2      3      IT      13.24      False      Excellent
3      4      Finance      1.55      False      Excellent
4      5      Marketing      4.16      True      Average
5      6      Finance      7.42      True      Excellent
6      7      Marketing      10.42      True      Good
7      8      Marketing      12.90      True      Good
8      9      Marketing      10.00      True      Poor
9     10      HR      18.74      False      Average
```

```
Salary First_Name Last_Name
0  54550.0   Michael   Davis
1  55225.0    Karen   Brown
2  83100.0   Joseph   Johnson
3  53875.0    David   Garcia
4  60400.0    Linda   Martinez
5  68550.0   Michael   Brown
6  76050.0   Charles   Moore
7  82250.0    David   Lopez
8  75000.0    David   Johnson
9  96850.0  Patricia   Martinez
```

We can also pass in a list of rows.

```
[19]: employee.iloc[[1,2,39]]
```

```
[19]: Employee_ID Department Years_of_Experience Full_Time Performance_Score \
1      2      IT      2.09      False      Poor
2      3      IT      13.24      False      Excellent
39     40      HR      11.64      True      Excellent
```

```
Salary First_Name Last_Name
1  55225.0    Karen   Brown
2  83100.0   Joseph   Johnson
39 79100.0    Karen   Martinez
```

This is particularly useful if we want rows that are not in a sequence. For example if you want rows indexed 0,2,39.

1.8.1 Indexing both axes

You can mix the indexer types for the index and columns. With scalar integers. For example:

```
[20]: employee.iloc[0,0]
```

```
[20]: 1
```

In the above code we wanted to look at row index 0 and column index 0. In the example below are looking at row index 0 and column index 1.

```
[21]: employee.iloc[0,1]
```

```
[21]: 'Finance'
```

With lists of integers. Remember the first list is for row and second list is for column.

```
[22]: employee.iloc[[0,1,2],[0,1,2]]
```

```
[22]:
```

	Employee_ID	Department	Years_of_Experience
0	1	Finance	1.82
1	2	IT	2.09
2	3	IT	13.24

We could also use range

```
[23]: employee.iloc[range(0,3),range(0,3)]
```

```
[23]:
```

	Employee_ID	Department	Years_of_Experience
0	1	Finance	1.82
1	2	IT	2.09
2	3	IT	13.24

With slice objects.

```
[24]: employee.iloc[0:3,0:4]
```

```
[24]:
```

	Employee_ID	Department	Years_of_Experience	Full_Time
0	1	Finance	1.82	False
1	2	IT	2.09	False
2	3	IT	13.24	False

The code selects the first three rows and the first four columns from the DataFrame named `employee` using integer-location based indexing.

1.9 Label-based selection— `.loc`

Access a group of rows and columns by label(s) or a boolean array. Single label. Note this returns the row as a Series. To see this clearly let's create a fake dataset

```
[25]: df1=pd.DataFrame([[1,2,2],[4,5,2],[7,5,2],[232,21,24]],  
                      index=['cobra','viper','sidewinder','rattle_snake'],  
                      columns=['max_speed','shield','windy']  
)
```

```
[26]: df1
```

```
[26]:
```

	max_speed	shield	windy
cobra	1	2	2
viper	4	5	2
sidewinder	7	5	2
rattle_snake	232	21	24

Single label. Note this returns the row as a Series.

```
[27]: df1.loc['viper']
```

```
[27]: max_speed    4
      shield      5
      windy       2
      Name: viper, dtype: int64
```

The code above selects the row(s) with the index label 'viper' from the DataFrame df1 using label-based indexing.

```
[28]: df1.loc[['viper', 'rattle_snake']]
```

```
[28]:
```

	max_speed	shield	windy
viper	4	5	2
rattle_snake	232	21	24

The code above selects the rows with index labels 'viper' and 'rattle_snake' from the DataFrame df1 using label-based indexing. In this case, the labels provided in the square brackets(i.e. list) ['viper', 'rattle_snake'] represent the index labels of the rows that you want to select from df1.

```
[29]: df1.loc['viper':'rattle_snake']
```

```
[29]:
```

	max_speed	shield	windy
viper	4	5	2
sidewinder	7	5	2
rattle_snake	232	21	24

The code above is used to select and retrieve a range of rows from a Pandas DataFrame named df1 using label-based indexing with the .loc[] method. By specifying the range 'viper':'rattle_snake', it tells Pandas to include all rows starting from the row labeled 'viper' up to and including the row labeled 'rattle_snake'. This operation will return all rows within this range, along with all their columns, effectively slicing the DataFrame based on row labels.

```
[30]: df1.loc['viper':'rattle_snake', 'shield':'windy']
```

```
[30]:
```

	shield	windy
viper	5	2
sidewinder	5	2
rattle_snake	21	24

The code above selects a subset of rows and columns from `df1` using label-based indexing with the `.loc[]` method. Specifically, it retrieves rows from ‘viper’ to ‘rattle_snake’ and columns from ‘shield’ to ‘windy’. This means it includes all rows starting at the index labeled ‘viper’ up to and including the row labeled ‘rattle_snake’, and similarly, it includes all columns starting at ‘shield’ up to and including ‘windy’. The result is a DataFrame that contains the specified slice of rows and columns, based on their labels.

1.10 Difference between `.loc` and `.iloc`.

`.loc[]` and `.iloc[]` are two indexing methods available in Pandas for selecting data from a DataFrame, but they cater to different types of indexing: label-based and integer-location based, respectively. The main difference between `.loc[]` and `.iloc[]` lies in how they interpret their arguments. `.loc[]` is used for label-based indexing, meaning it selects data based on data frame index or column names. For example, `df.loc['row_label', 'column_label']` retrieves the data at the specified row and column labels. On the other hand, `.iloc[]` is used for integer-location based indexing, meaning it selects data based on the integer positions of the rows and columns. So, `df.iloc[1, 2]` would retrieve the data located at the second row and third column, as indexing starts at 0. While `.loc[]` allows for more human-readable code by using explicit labels, `.iloc[]` offers a straightforward way to navigate through data based on numeric positions, making it particularly useful when the row or column labels are not known or when iterating through data sequentially.

[]: