

Untitled4

March 21, 2024

1 Week 7 Python

```
[1]: import pandas as pd
```

1.1 Grouping, Aggregating and Filtering

Let's import employee_datasetv2.csv dataset

```
[2]: df=pd.read_csv('employee_datasetv2.csv')
```

```
[3]: df.head(10)
```

```
[3]:
```

	Employee_ID	Department	Years_of_Experience	Full_Time	Performance_Score	\
0	1	Finance	1.820000	False	Good	
1	2	IT	2.090000	False	Poor	
2	3	IT	7.047228	False	Excellent	
3	4	Finance	1.237509	False	Excellent	
4	5	Marketing	4.160000	True	Average	
5	6	Finance	2.475017	True	Excellent	
6	7	Marketing	10.420000	True	Good	
7	8	Marketing	6.102669	True	Good	
8	9	Marketing	10.000000	True	Poor	
9	10	HR	7.479808	False	Average	

	Salary	First_Name	Last_Name	Date_of_Birth	Gender	...	\
0	56550.0	Michael	Davis	2000-05-11 05:33:39	Other	...	
1	55225.0	Karen	Brown	1984-12-18 14:13:34	Other	...	
2	87600.0	Joseph	Johnson	1991-07-29 11:32:58	Male	...	
3	53875.0	David	Garcia	1998-01-24 12:21:27	Other	...	
4	61400.0	Linda	Martinez	1973-06-05 08:09:32	Other	...	
5	70550.0	Michael	Brown	1973-03-07 10:44:36	Female	...	
6	82050.0	Charles	Moore	1981-03-26 17:29:10	Male	...	
7	84750.0	David	Lopez	1981-03-25 12:42:58	Female	...	
8	80500.0	David	Johnson	1973-08-09 18:43:34	Female	...	
9	101350.0	Patricia	Martinez	1967-11-27 08:26:50	Female	...	

	Location	Remote_Work	Hire_Date	\
--	----------	-------------	-----------	---

0	Glasgow, Scotland	True	2022-05-16 07:30:27
1	London, England	False	2020-05-03 22:30:19
2	Belfast, Northern Ireland	True	2017-02-25 23:43:54
3	Belfast, Northern Ireland	True	2022-12-18 10:59:47
4	Manchester, England	True	2012-04-04 14:02:58
5	Glasgow, Scotland	False	2021-09-22 12:23:53
6	Glasgow, Scotland	False	2011-08-25 03:34:26
7	Belfast, Northern Ireland	False	2018-02-05 19:31:10
8	Glasgow, Scotland	False	2012-09-29 19:05:03
9	Cardiff, Wales	True	2016-09-20 00:26:27

	Last_Performance_Review_Date	Project_Count	Bonus	\
0	2022-10-08 07:30:27	5	0	
1	2020-08-15 22:30:19	6	500	
2	2017-07-17 23:43:54	3	2000	
3	2023-04-30 10:59:47	9	500	
4	2013-01-12 14:02:58	5	500	
5	2022-06-05 12:23:53	8	0	
6	2012-01-13 03:34:26	1	1500	
7	2018-06-13 19:31:10	5	500	
8	2013-03-30 19:05:03	3	1000	
9	2017-05-01 00:26:27	1	1500	

	Employee_Satisfaction_Score	Department_Budget	Training_Hours_Last_Year	\
0	2.68	75000	46	
1	1.99	100000	15	
2	2.42	50000	4	
3	4.03	125000	34	
4	1.06	75000	11	
5	1.46	150000	24	
6	1.18	75000	20	
7	1.16	150000	35	
8	4.42	100000	22	
9	3.81	150000	15	

	Number_of_Direct_Reports
0	6
1	8
2	3
3	3
4	5
5	2
6	5
7	6
8	9
9	9

[10 rows x 21 columns]

```
[4]: df.describe()
```

```
[4]:
```

	Employee_ID	Years_of_Experience	Salary	Project_Count	\
count	50.00000	50.000000	50.000000	50.000000	
mean	25.50000	6.552038	78774.000000	4.900000	
std	14.57738	3.777367	15041.086553	2.697315	
min	1.00000	1.120000	53875.000000	1.000000	
25%	13.25000	2.980262	65975.000000	2.250000	
50%	25.50000	6.184805	80150.000000	5.000000	
75%	37.75000	10.099589	89156.250000	7.000000	
max	50.00000	13.880903	105750.000000	9.000000	

	Bonus	Employee_Satisfaction_Score	Department_Budget	\
count	50.000000	50.000000	50.000000	
mean	1090.000000	2.684200	98000.000000	
std	760.571751	1.188431	38412.370456	
min	0.000000	1.060000	50000.000000	
25%	500.000000	1.415000	50000.000000	
50%	1250.000000	2.710000	100000.000000	
75%	2000.000000	3.562500	125000.000000	
max	2000.000000	4.840000	150000.000000	

	Training_Hours_Last_Year	Number_of_Direct_Reports
count	50.000000	50.000000
mean	24.840000	4.620000
std	14.945288	2.834842
min	0.000000	0.000000
25%	11.250000	2.000000
50%	23.000000	5.000000
75%	38.000000	7.000000
max	49.000000	9.000000

1.2 Grouping

Concept: Grouping involves organizing data into groups based on some criteria, such as department or education level. Pandas uses the `.groupby()` method to achieve this, allowing for powerful and flexible data analysis.

Example: Group employees by their department to calculate the average salary within each department.

```
[5]: grouped_by_dept=df.groupby('Department')
```

This line of code performs a grouping operation on the DataFrame `df` based on the values in the 'Department' column. Here's a step-by-step explanation:

1. **df:** This is the DataFrame that contains the data from the `employee_datasetv2.csv` file

you've read into Python using **pandas**. It's assumed to have multiple columns, one of which is named 'Department'.

2. **.groupby('Department')**: This method is called on the DataFrame **df**. The **groupby** method is used to split the data into groups based on some criteria. In this case, the criteria is the 'Department' column. This means that the DataFrame will be divided into groups where each group contains rows that have the same value in the 'Department' column.
3. **grouped_by_dept**: This is the variable to which the result of the **groupby** operation is assigned. However, it's important to note that **grouped_by_dept** does not hold a simple DataFrame. Instead, it holds a **DataFrameGroupBy** object, which is a collection of groups awaiting further analysis or manipulation.

After this line of code is executed, **grouped_by_dept** can be used to perform operations on each group separately, such as calculating statistics (mean, median, max, min, etc.), applying functions, or aggregating data in various ways specific to each department. This is particularly useful for analyzing differences across departments or summarizing data at the department level.

```
[6]: avg_salary_by_dept=grouped_by_dept['Salary'].mean()
```

This line of code calculates the average salary within each department from the previously grouped DataFrame (**grouped_by_dept**). Here's how it works:

1. **grouped_by_dept**: This is the **DataFrameGroupBy** object you created by grouping the original DataFrame **df** by the 'Department' column. It represents your dataset partitioned into groups, where each group consists of rows that have the same department name.
2. **['Salary']**: This part of the code selects the 'Salary' column from each group. Because **grouped_by_dept** is a group of data grouped by department, this operation is applied to each department group, effectively isolating the 'Salary' data for further computation within each department.
3. **.mean()**: This method computes the mean (average) of the 'Salary' column for each department group. The operation is performed on the 'Salary' column of each department, calculating the average salary for that department.
4. **avg_salary_by_dept**: The result of the mean calculation for each group is assigned to this variable. The resulting object is a Series where the index is the department names (from the 'Department' column used to group the DataFrame) and the values are the calculated average salaries for those departments.

The final outcome is that **avg_salary_by_dept** holds the average salary for each department, allowing for an easy comparison of how average salaries differ across departments in the dataset.

```
[7]: avg_salary_by_dept
```

```
[7]: Department
Finance      76141.666667
HR           78359.090909
IT           80515.625000
Marketing    80796.875000
Name: Salary, dtype: float64
```

```
[8]: avg_salary_by_dept.reset_index()
```

```
[8]:   Department      Salary
0    Finance  76141.666667
1         HR  78359.090909
2         IT  80515.625000
3  Marketing  80796.875000
```

The `avg_salary_by_dept.reset_index()` method is used to transform the Series `avg_salary_by_dept` into a DataFrame and reset its index. This operation is particularly useful after performing groupby calculations which often result in the grouping columns becoming the index of the resulting Series or DataFrame. Here's the detailed process:

1. **avg_salary_by_dept:** Before calling `reset_index()`, `avg_salary_by_dept` is a pandas Series with departments as its index and their corresponding average salaries as its values. This Series was obtained by grouping the original DataFrame `df` by 'Department', selecting the 'Salary' column, and then calculating the mean salary for each department.
2. **.reset_index():** This method converts the Series into a DataFrame and resets its index. By default, the index of the Series (in this case, the department names) becomes a regular column in the resulting DataFrame, and a new numerical index is introduced, starting from 0 and incrementing by 1 for each row.
3. **Result:** The result of this operation is a DataFrame with two columns:
 - The first column (often named 'index' if no other arguments are provided to `reset_index()`) contains the department names that were previously the index in the Series.
 - The second column contains the average salaries for each department. The name of this column will be the same as the name of the Series, if it had one, or it might default to a generic name like '0' if the Series did not have a name.

This process makes the data structure more flexible for further analysis or for merging with other DataFrames, as it converts the Series back into a DataFrame and provides a standard, numerical index.

1.3 Grouping and Aggregating

Concept: Aggregation refers to any data transformation that produces scalar values from arrays. After grouping data, you can perform aggregation operations like calculating sums, means, and minimums or maximums.

Example: Find the total bonuses distributed in each department.

```
[9]: total_bonus_by_dept= df.groupby('Department')['Bonus'].sum()
```

This line of code calculates the total sum of bonuses within each department from your DataFrame `df`. It's doing several things in a single, concise operation:

1. **df.groupby('Department'):** This groups the DataFrame `df` by the 'Department' column. Each group consists of all rows that have the same value in the 'Department' column, effectively segregating the data by department.

2. `['Bonus']`: After grouping, this selects the 'Bonus' column from each group. This means that for each department, only the 'Bonus' data is considered for the subsequent operation.
3. `.sum()`: This method is applied to the 'Bonus' column of each department group to calculate the sum of bonuses within that department. It aggregates the bonus values by adding them up for each department.
4. `total_bonus_by_dept`: The result of the sum operation is assigned to this variable. The resulting object is a pandas Series where the index is the department names (as determined by the 'Department' column used to group the DataFrame) and the values are the total sum of bonuses for those departments.

Thus, `total_bonus_by_dept` holds the total bonus amount for each department, enabling an analysis of how bonuses are distributed across different departments within the organization. This can be useful for understanding departmental rewards or for budgeting and financial planning related to employee compensation.

```
[10]: total_bonus_by_dept.reset_index()
```

```
[10]:  Department  Bonus
0    Finance  14500
1         HR   12500
2         IT   20000
3  Marketing   7500
```

The `total_bonus_by_dept.reset_index()` method is used to convert the `total_bonus_by_dept` Series into a DataFrame and reset its index. This transformation is especially useful when you have a Series with a non-default index (in this case, department names as the index) and you want to turn it back into a column, making the data structure resemble a traditional table with a default integer index. Here's what happens during this operation:

1. `total_bonus_by_dept`: Initially, this is a pandas Series resulting from summing up the 'Bonus' values for each department in the DataFrame `df`. The index of this Series is the unique values from the 'Department' column, representing different departments.
2. `.reset_index()`: This method performs two main actions:
 - It converts the Series into a DataFrame. The values of the Series (total bonuses) become one column in the new DataFrame.
 - It resets the index of this DataFrame to a default integer index, starting from 0. The original index (department names) is added as a new column in the DataFrame.
3. **Result**: The outcome is a DataFrame with two columns:
 - The first column, often named 'Department' if the original Series had a name or defaulting to 'index' if not specified, contains the names of the departments. These were previously the index of the Series.
 - The second column contains the corresponding total bonus amounts for each department. If the Series had a name (in this case, 'Bonus'), that name would be used for this column. Otherwise, a default name would be applied, typically a numerical label like '0'.

This operation is particularly useful for preparing the data for further analysis, reporting, or visualization, as it presents the information in a more conventional tabular format with a simple numerical index.

1.4 Combining Grouping and Aggregation

Concept: By **combining grouping and aggregation**, you can perform more complex analyses that involve summarizing grouped data using various aggregation functions.

Example: Calculate the average number of training hours last year, grouped by department and full-time/part-time status.

```
[11]: avg_training_hours=df.  
      ↳groupby(['Department','Full_Time'])['Training_Hours_Last_Year'].mean()
```

This line of code calculates the average training hours in the last year for each department, further segmented by whether the employee is full-time or not. Here's a breakdown of what each part does:

1. `df.groupby(['Department','Full_Time'])`: This groups the DataFrame `df` by two columns: 'Department' and 'Full_Time'. It creates groups based on the unique combinations of department names and full-time status. This means that for each department, there will be separate groups for full-time employees and for those who are not full-time, allowing for a more nuanced analysis.
2. `['Training_Hours_Last_Year']`: From each group created in the step above, this selects the 'Training_Hours_Last_Year' column. This column presumably contains numerical values representing the number of training hours each employee completed in the last year.
3. `.mean()`: This computes the mean (average) of the 'Training_Hours_Last_Year' values within each group. Since the groups are defined by both department and full-time status, this calculates the average training hours for each category of employee within each department.
4. `avg_training_hours`: The result of the mean calculation is assigned to this variable. The resulting object is likely a pandas Series with a MultiIndex. The first level of the index is the 'Department', and the second level is the 'Full_Time' status (which could be a boolean or some other indicator of whether an employee is considered full-time). The values of this Series are the average training hours for each combination of department and full-time status.

This operation allows you to understand not just how training hours are distributed across different departments, but also how this distribution might vary between full-time and part-time (or equivalent distinctions) employees within those departments. It's a powerful way to analyze the data for insights into training practices across different segments of the workforce.

```
[12]: avg_training_hours
```

```
[12]: Department  Full_Time  
Finance      False      30.000000  
              True       27.888889  
HR            False      22.000000  
              True       25.285714
```

```

IT          False      21.142857
           True       23.222222
Marketing   False      26.500000
           True       22.666667
Name: Training_Hours_Last_Year, dtype: float64

```

To illustrate the differences between `avg_training_hours.reset_index()` and `avg_training_hours.unstack()`, let's first clarify what each method does, especially in the context of a pandas Series with a MultiIndex created from grouping by multiple columns, as is the case with `avg_training_hours`.

1.4.1 `avg_training_hours.reset_index()`

- **What it does:** Converts the Series with a MultiIndex (in this case, 'Department' and 'Full_Time') into a DataFrame. The indices ('Department' and 'Full_Time') become columns in the resulting DataFrame, making it a "flat" structure with a default integer index.
- **Use case:** Useful when you want a simple DataFrame format where each group's identifying information ('Department' and 'Full_Time') is moved to columns alongside the values (average training hours).

1.4.2 `avg_training_hours.unstack()`

- **What it does:** Pivots the level of the specified index columns (if no level is specified, the last level is unstacked) to the columns, creating a new DataFrame where each unique value of the unstacked level becomes a column, and the values are the data points corresponding to each combination of indices.
- **Use case:** Useful for creating a pivot table-like structure where one of the indices ('Department' or 'Full_Time') is used to create columns, making it easier to compare across one of the categorical dimensions.

To understand the practical difference, let's apply both methods to the `avg_training_hours` Series. This will give us a direct comparison of the outputs.

Let's proceed by loading the data, performing the operations to create `avg_training_hours`, and then applying both `reset_index()` and `unstack()` to see the differences.

```
[13]: avg_training_hours.reset_index()
```

```

[13]:  Department  Full_Time  Training_Hours_Last_Year
0      Finance    False      30.000000
1      Finance     True      27.888889
2         HR     False      22.000000
3         HR     True      25.285714
4         IT     False      21.142857
5         IT     True      23.222222
6  Marketing    False      26.500000
7  Marketing     True      22.666667

```

```
[14]: avg_training_hours.unstack()
```



```
[14]: Full_Time      False      True
      Department
Finance      30.000000  27.888889
HR           22.000000  25.285714
IT           21.142857  23.222222
Marketing    26.500000  22.666667
```

Based on the operations performed, here's a comparison of the results:

1.4.3 Using `reset_index()`

The `reset_index()` method transformed `avg_training_hours` into a DataFrame with three columns: 'Department', 'Full_Time', and 'Training_Hours_Last_Year'. Each row represents a unique combination of 'Department' and 'Full_Time', with the corresponding average training hours listed. For example, the Finance department has separate rows for full-time (True) and not full-time (False) employees, with average training hours of approximately 27.89 and 30.00 hours, respectively.

1.4.4 Using `unstack()`

The `unstack()` method, on the other hand, pivoted the 'Full_Time' level of the index to the columns, creating a DataFrame where each row represents a department, and there are separate columns for full-time (True) and not full-time (False) employees' average training hours. This structure makes it easy to compare the average training hours between full-time and not full-time employees within each department at a glance. For instance, in the Finance department, full-time employees have an average of approximately 27.89 training hours, while not full-time employees have an average of 30.00 hours.

Both methods are useful for data analysis, but they serve different purposes depending on your needs. `reset_index()` gives a "flattened" DataFrame that can be useful for further analysis or merging with other data, while `unstack()` provides a pivot table-like structure that is great for comparing subgroups side by side.

1.5 Apply Method

Concept: Pandas also allows for more advanced group operations such as filtering, transforming, and **applying** custom functions to groups.

Scenario: Suppose we want to categorize employees based on their years of experience into 'Junior', 'Mid-Level', and 'Senior'.

```
[15]: def cat_exp(years):
      if years<3:
          return 'Junior'
      elif years<7:
          return 'Mid-Level'
      else:
          return 'Senior'
```

```
[16]: df['Experience_Level']=df['Years_of_Experience'].apply(cat_exp)
```

This block of code defines a function named `cat_exp` that categorizes years of experience into three levels: 'Junior', 'Mid-Level', and 'Senior'. It then applies this function to the 'Years_of_Experience' column of the DataFrame `df` to create a new column named 'Experience_Level' with the categorized experience levels. Here's a step-by-step explanation:

1. **Define the `cat_exp` function:**

- The function takes a single argument `years`, which represents the number of years of experience.
- It contains a series of conditional statements (`if`, `elif`, `else`) to categorize the experience based on the number of years:
 - If `years` is less than 3, it returns 'Junior'.
 - If `years` is 3 or more but less than 7, it returns 'Mid-Level'.
 - Otherwise, it returns 'Senior'.

2. **Apply the function to the DataFrame:**

- `df['Years_of_Experience'].apply(cat_exp)` applies the `cat_exp` function to each element in the 'Years_of_Experience' column of `df`. The `.apply()` method is used to apply a function along an axis of the DataFrame or on values of a Series. Here, it's used on a Series to transform each year value into a categorical label according to the logic defined in `cat_exp`.

3. **Create a new column 'Experience_Level':**

- The result of the `.apply()` operation, which is a Series of categorized experience levels, is assigned to a new column in `df` named 'Experience_Level'. This operation adds the 'Experience_Level' column to the DataFrame, where each row's value corresponds to the categorized experience level of the respective 'Years_of_Experience'.

By executing this code, `df` now includes a new column 'Experience_Level' that categorizes each employee's experience based on the number of years they have worked, providing a straightforward way to segment the workforce into experience-based groups.

```
[17]: df.head(10)
```

```
[17]: Employee_ID Department Years_of_Experience Full_Time Performance_Score \
0          1      Finance          1.820000         False           Good
1          2          IT          2.090000         False           Poor
2          3          IT          7.047228         False      Excellent
3          4      Finance          1.237509         False      Excellent
4          5  Marketing          4.160000          True       Average
5          6      Finance          2.475017          True      Excellent
6          7  Marketing         10.420000          True           Good
7          8  Marketing          6.102669          True           Good
8          9  Marketing         10.000000          True           Poor
9         10          HR          7.479808         False       Average

      Salary First_Name Last_Name      Date_of_Birth Gender  ...  \
0    56550.0    Michael    Davis  2000-05-11 05:33:39  Other  ...
1    55225.0      Karen    Brown  1984-12-18 14:13:34  Other  ...
```

2	87600.0	Joseph	Johnson	1991-07-29	11:32:58	Male	...
3	53875.0	David	Garcia	1998-01-24	12:21:27	Other	...
4	61400.0	Linda	Martinez	1973-06-05	08:09:32	Other	...
5	70550.0	Michael	Brown	1973-03-07	10:44:36	Female	...
6	82050.0	Charles	Moore	1981-03-26	17:29:10	Male	...
7	84750.0	David	Lopez	1981-03-25	12:42:58	Female	...
8	80500.0	David	Johnson	1973-08-09	18:43:34	Female	...
9	101350.0	Patricia	Martinez	1967-11-27	08:26:50	Female	...

	Remote_Work	Hire_Date	Last_Performance_Review_Date	\
0	True	2022-05-16 07:30:27	2022-10-08 07:30:27	
1	False	2020-05-03 22:30:19	2020-08-15 22:30:19	
2	True	2017-02-25 23:43:54	2017-07-17 23:43:54	
3	True	2022-12-18 10:59:47	2023-04-30 10:59:47	
4	True	2012-04-04 14:02:58	2013-01-12 14:02:58	
5	False	2021-09-22 12:23:53	2022-06-05 12:23:53	
6	False	2011-08-25 03:34:26	2012-01-13 03:34:26	
7	False	2018-02-05 19:31:10	2018-06-13 19:31:10	
8	False	2012-09-29 19:05:03	2013-03-30 19:05:03	
9	True	2016-09-20 00:26:27	2017-05-01 00:26:27	

	Project_Count	Bonus	Employee_Satisfaction_Score	Department_Budget	\
0	5	0	2.68	75000	
1	6	500	1.99	100000	
2	3	2000	2.42	50000	
3	9	500	4.03	125000	
4	5	500	1.06	75000	
5	8	0	1.46	150000	
6	1	1500	1.18	75000	
7	5	500	1.16	150000	
8	3	1000	4.42	100000	
9	1	1500	3.81	150000	

	Training_Hours_Last_Year	Number_of_Direct_Reports	Experience_Level
0	46	6	Junior
1	15	8	Junior
2	4	3	Senior
3	34	3	Junior
4	11	5	Mid-Level
5	24	2	Junior
6	20	5	Senior
7	35	6	Mid-Level
8	22	9	Senior
9	15	9	Senior

[10 rows x 22 columns]

The `.apply()` method in pandas is a powerful and flexible tool used to apply a function along an axis of a DataFrame or to elements of a Series. It allows for both row-wise and column-wise operations in a DataFrame, and element-wise operations in a Series. Here's a more detailed explanation:

1.5.1 Applying to a DataFrame

When used on a DataFrame, `.apply()` can operate across rows or columns:

- **Column-wise operation (`axis=0`):** By default, or when `axis=0` is specified, `.apply()` applies the given function to each column, treating each column as an array-like structure. This is useful for performing operations that need to consider each value in a column, like summing values or finding maximum values.
- **Row-wise operation (`axis=1`):** When `axis=1` is specified, `.apply()` applies the function to each row. This mode is handy for operations that need to consider values across columns for each row, such as calculating a sum of certain columns for each row.

1.5.2 Applying to a Series

When `.apply()` is used on a Series, it applies a given function to each element in the Series. This is similar to mapping a function over an iterable, transforming each element in the Series according to the logic defined in the function.

1.5.3 General Use Cases

- **Data transformation:** You can use `.apply()` to perform complex data transformations, such as converting data types, applying conditional logic (like categorizing data based on values), or combining data from multiple columns.
- **Aggregation and summary statistics:** While pandas provides built-in methods for common operations (like `.sum()`, `.mean()`, etc.), `.apply()` allows for custom aggregation or summary statistics that aren't directly supported.
- **Applying custom functions:** `.apply()` shines when you have a specific, possibly complex operation that needs to be performed on a dataset, which is not covered by pandas' built-in methods. This includes applying mathematical formulas, data cleaning operations, or any user-defined function.

1.5.4 Important Considerations

- **Performance:** While `.apply()` is very flexible, it's not always the fastest option for large datasets, especially with complex operations. Vectorized operations using pandas' built-in methods or accessing the underlying NumPy arrays can be more efficient.
- **Broadcasting:** The function used with `.apply()` should return a value that pandas can incorporate into a Series or DataFrame. Depending on the operation, you might need to ensure that the return values are of a consistent format or type.

In summary, `.apply()` is a versatile method that can significantly enhance the data manipulation capabilities of pandas, allowing for custom, row-wise, column-wise, or element-wise operations on data structures.

1.6 Task 1:

calculate total salary: this should include the existing bonus; and added bonus for individuals that has a score of 4.00 or more. Added bounus is 1000 pounds

Here's how you can do it:

1. Define a function that takes an employee's satisfaction score as its input.
2. Inside the function, use an if statement to check if the score is 4.00 or higher.
3. Return 1000 if the condition is met, indicating the employee receives the added bonus.
4. Return 0 otherwise, indicating no added bonus for that employee.
5. Apply this function to the `Employee_Satisfaction_Score` column to create the `Added_Bonus` column.

Let's define the function and apply it:

```
[18]: # Define the function to determine the added bonus
def calculate_added_bonus(satisfaction_score):
    if satisfaction_score >= 4.00:
        return 1000
    else:
        return 0

# Apply the function to the Employee_Satisfaction_Score column
df['Added_Bonus'] = df['Employee_Satisfaction_Score'].
    ↪ apply(calculate_added_bonus)

# Recalculate the Total_Salary to reflect any changes
df['Total_Salary'] = df['Salary'] + df['Bonus'] + df['Added_Bonus']

# Display the first few rows to verify the calculation
df[['Employee_Satisfaction_Score', 'Added_Bonus', 'Total_Salary']].head()
```

```
[18]: Employee_Satisfaction_Score  Added_Bonus  Total_Salary
0                2.68                0      56550.0
1                1.99                0      55725.0
2                2.42                0      89600.0
3                4.03             1000      55375.0
4                1.06                0      61900.0
```

The function to determine the added bonus based on the `Employee_Satisfaction_Score` has been successfully defined and applied. The `Added_Bonus` column now reflects the additional 1000 pounds for employees with a satisfaction score of 4.00 or higher, as calculated by the custom function. The `Total_Salary` column has been recalculated to include this `Added_Bonus`. Here are the results for the first few rows:

- The `Employee_Satisfaction_Score` column shows the satisfaction scores of the employees.
- The `Added_Bonus` column indicates the additional bonus, where only employees with a score of 4.00 or more receive the 1000 pounds bonus. For instance, the employee at index 3, with a satisfaction score of 4.03, receives the added bonus, resulting in an `Added_Bonus` of 1000

pounds.

- The `Total_Salary` column has been updated to include any added bonuses alongside the original salary and bonus, accurately reflecting the total compensation for each employee.

1.7 Task 2

Split the `Location` column into `City` and `Country`. For example `Glasgow, Scotland` should be split into two `Glasgow` and `Scotland`

To split the `Location` column in your `DataFrame` into two separate columns, `City` and `Country`, you can use the `.str.split()` method on the `Location` column with the appropriate separator. Assuming the format of the `Location` values is consistent and uses a comma as the separator (as in “Glasgow, Scotland”), here’s how you can do it:

1. Use the `.str.split()` method on the `Location` column, specifying `,` as the separator and `expand=True` to return a `DataFrame`.
2. Assign the result to two new columns in your existing `DataFrame`, `City` and `Country`. Let’s implement this approach.

```
[19]: # Split the Location column into City and Country
df[['City', 'Country']] = df['Location'].str.split(',', expand=True)

# Display the first few rows to verify the split
df[['Location', 'City', 'Country']].head()
```

```
[19]:
```

	Location	City	Country
0	Glasgow, Scotland	Glasgow	Scotland
1	London, England	London	England
2	Belfast, Northern Ireland	Belfast	Northern Ireland
3	Belfast, Northern Ireland	Belfast	Northern Ireland
4	Manchester, England	Manchester	England

The `Location` column has been successfully split into two new columns, `City` and `Country`. For example, “Glasgow, Scotland” has been split into “Glasgow” for the city and “Scotland” for the country, as seen in the first row. This process has been applied across the `DataFrame`, allowing for more specific geographic analysis or filtering based on either city or country.

```
[ ]:
```